



Deep Learning From Scratch

# TensorFlow

Santi Seguí

# TensorFlow: Just another library for Deep Learning?



theano

Caffe

# What is TensorFlow

- TensorFlow is a deep learning library recently open-sourced by Google.
- But, what does it actually do?
  - Provides primitives for defining **functions on tensors** and **automatically computing their derivatives**

# What is a **Tensor**?

A typed multi-dimensional array

# Simple **Numpy** Recap

```
In [2]: # Simple Numpy Recap  
import numpy as np  
  
a = np.zeros((2,2)); b = np.ones((2,2))  
print np.sum(b,axis=1)  
  
[ 2.  2.]
```

```
In [3]: print a.shape  
  
(2, 2)
```

```
In [4]: print np.reshape(a,(1,4))  
  
[[ 0.  0.  0.  0.]]
```

# Repeat in TensorFlow

```
In [5]: import tensorflow as tf
```

```
In [6]: sess = tf.InteractiveSession()
```

What is **InteractiveSession**?

```
In [7]: a = tf.zeros((2,2)); b = tf.ones((2,2))
```

```
In [8]: tf.reduce_sum(b, reduction_indices=1).eval()
```

eval()?

```
Out[8]: array([ 2.,  2.], dtype=float32)
```

```
In [9]: a.get_shape()
```

```
Out[9]: TensorShape([Dimension(2), Dimension(2)])
```

```
In [10]: tf.reshape(a, (1, 4)).eval()
```

```
Out[10]: array([[ 0.,  0.,  0.,  0.]], dtype=float32)
```

```
In [11]: a = np.zeros((2,2))
         ta = tf.zeros((2,2))
         print a
         print ta
```

```
[[ 0.  0.]
 [ 0.  0.]]
Tensor("zeros_1:0", shape=(2, 2), dtype=float32)
```

TensorFlow computations define a **computational graph** that has not a numerical value until explicit evaluation.

# TensorFlow Mechanics

1. Prepare the Data

1. Inputs and Placeholders

2. Build the Graph

1. Inference

2. Loss

3. Training

3. Train The model

1. The Session

2. Compute Graph ops

3. Train loop

4. Evaluate the model

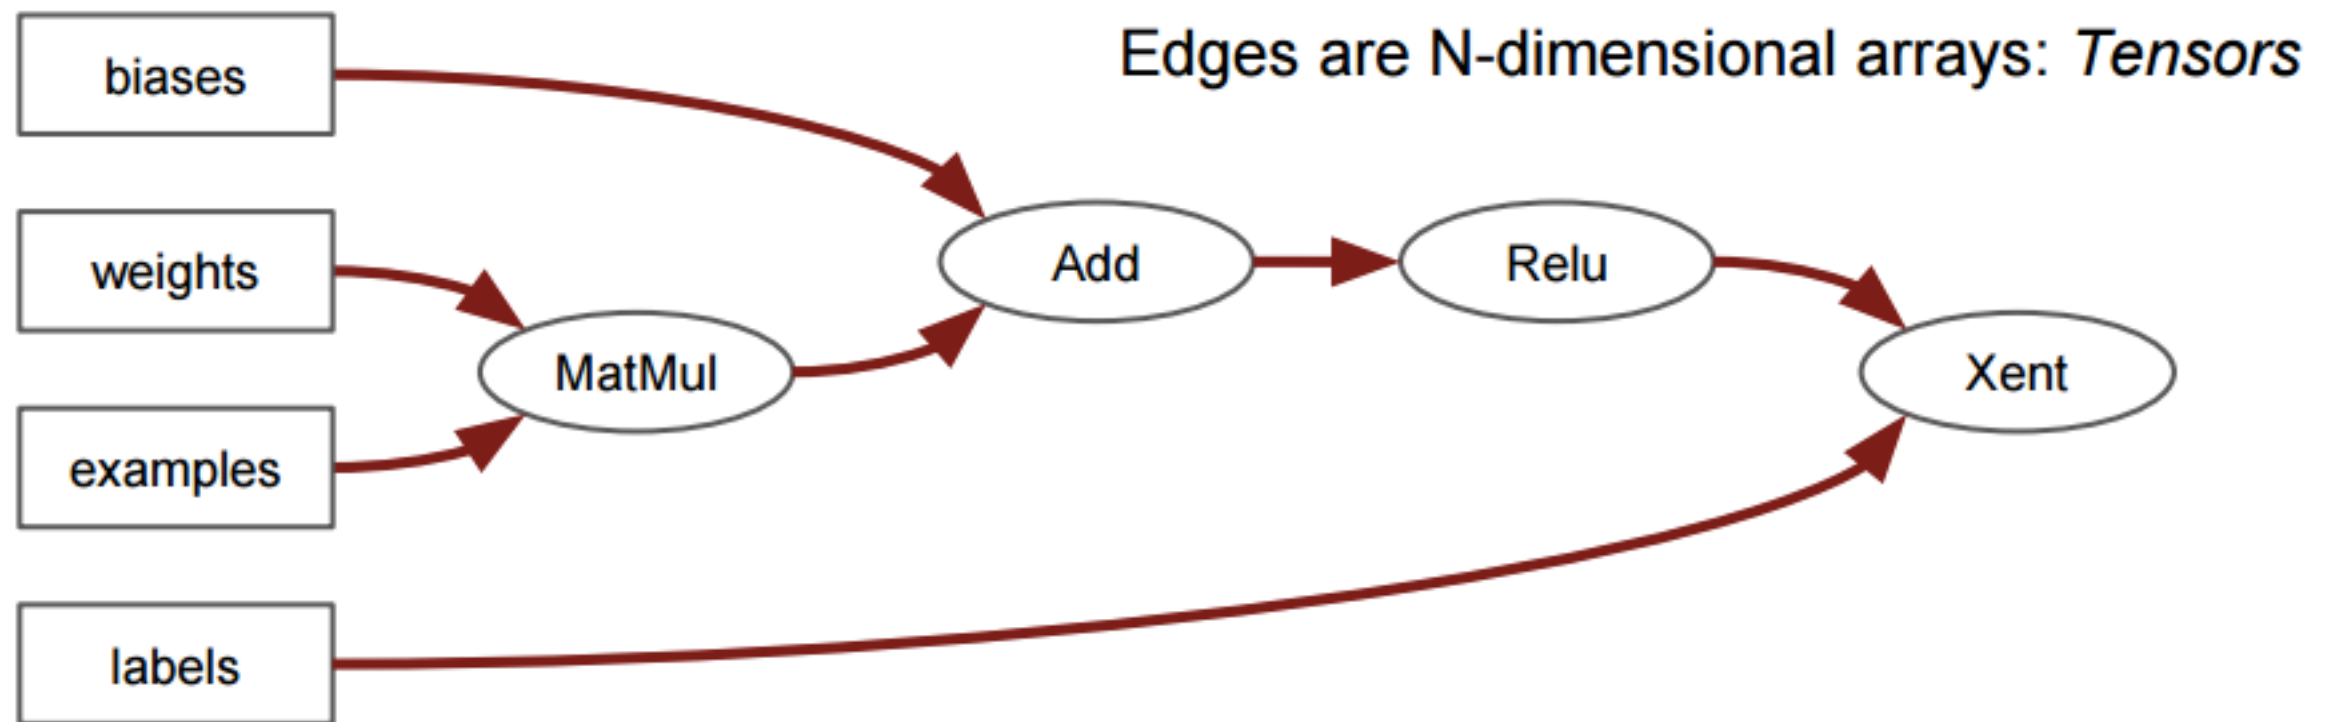


# TensorFlow Graph

- “TensorFlow programs are usually structured into a **construction phase**, that **assembles** a **graph**, and an **execution phase** that uses a session to **execute ops** in the **graph**”
- All **computations** add **nodes** to global default **graph**



# TensorFlow Graph



# TensorFlow Session

A Session object **encapsulates the environment** in which

Operation objects are executed, and Tensor objects are evaluated. A session may **own resources**, such as variables, queues, and readers. It is important to release these resources when they are no longer required.

Different ways to use TensorFlow sessions:

1) Using the Session object:

```
a = tf.constant(5.0)
b = tf.constant(6.0)
c = a * b
sess = tf.Session()
print sess.run(c)
sess.close()
```

2) Using the context manager:

```
a = tf.constant(5.0)
b = tf.constant(6.0)
c = a * b
with tf.Session() as sess:
    print(c.eval())
```

3) Using Interactive Session:

```
sess = tf.InteractiveSession()
a = tf.constant(5.0)
b = tf.constant(6.0)
c = a * b
print(c.eval())
sess.close()
```

# Feeding Data

- TensorFlow's **feed** mechanism lets you inject **data** into any Tensor in a **computation graph**.
- While you can replace any Tensor with feed data, including variables and constants, the best practice is to use a **placeholder** op node. A placeholder exists solely to serve as the target of feeds. It is not initialized and contains no data
- A **feed\_dict** is a python dictionary mapping from tf.placeholder vars (or their names) to data (numpy arrays, lists, etc.).

# Feeding Data

## Placeholders and feed dictionaries

```
In [1]: import tensorflow as tf
```

```
In [2]: input1 = tf.placeholder(tf.float32)
input2 = tf.placeholder(tf.float32)
```

```
In [3]: output = tf.mul(input1,input2)
```


```
In [4]: print output
```

```
Tensor("Mul:0", dtype=float32)
```

```
In [5]: with tf.Session() as sess:
        print(sess.run([output], feed_dict={input1:[7.],input2:[3.]}))
```

```
[array([ 21.], dtype=float32)]
```

# Feeding Data



# TensorFlow Variables

- “When you train a model, you use variables to hold and update parameters. Variables are in-memory buffers containing tensors. “
- “They must be explicitly initialized and can be saved to disk during and after training. You can later restore saved values to exercise or analyze the model.”

# TensorFlow Variables

```
In [5]: W = tf.Variable(tf.zeros([784, 10]))  
b = tf.Variable(tf.zeros([10]))
```

Why zeros?

```
In [6]: W = tf.Variable(tf.random_normal([784, 10], stddev=0.35), name = "weights")  
b = tf.Variable(tf.random_normal([10], stddev=0.35), name= "biases")
```

Variable initializers must be run explicitly before other ops in your model can be run. The easiest way to do that is to add an op that runs all the variable initializers, and run that op before using the model.

```
...  
# Add an op to initialize the variables.  
init_op = tf.initialize_all_variables()  
  
# Later, when launching the model  
with tf.Session() as sess:  
    #Run the init operation.  
    sess.run(init_op)  
    ...  
    # Use the model  
    ...
```



# Loss Functions

- The `loss()` function further builds the graph by adding the required loss ops.
- The cost function to be minimized during training can be specified easily.

# Train the Model

- Now that we have defined our **model** and **training cost function**, it is straightforward to train using TensorFlow. Because TensorFlow knows the entire computation graph, it can use **automatic differentiation** to find the gradients of the cost with respect to each of the variables.
- TensorFlow has a variety of builtin optimization algorithms.

# Optimization Functions

- AdamOptimizer
- GradientDescentOptimizer
- AdagradOptimizer
- AdadeltaOptimizer
- MomentumOptimizer
- FtrlOptimizer
- RMSPropOptimizer


# Linear Regression Example

```
In [3]: # Load data. Advertising dataset from "An Introduction to Statistical Learning",
# textbook by Gareth James, Robert Tibshirani, and Trevor Hastie
import numpy as np
data = pd.read_csv('dataset/Advertising.csv', index_col=0)
train_X = data[['TV']].values

train_Y = data.Sales.values
train_Y = train_Y[:, np.newaxis]

n_samples = train_X.shape[0]
print n_samples
print train_X.shape, train_Y.shape
plt.plot(train_X, train_Y, 'ro', label='Original data')
plt.show()
```

```
200
(200, 1) (200, 1)
```



# Linear Regression Example

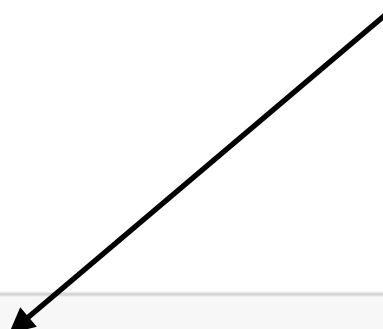
```
In [4]: import tensorflow as tf
# Define tf Graph Inputs
X = tf.placeholder("float", [None, 1])
y = tf.placeholder("float", [None, 1])

# Create Model variables
# Set model weights
W = tf.Variable(np.random.randn(), name="weight")
b = tf.Variable(np.random.randn(), name="bias")

# Construct a linear model
y_pred = tf.add(tf.mul(X, W), b)
```

# Linear Regression Example

$$J(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - (Wx_i + b))^2$$



```
In [ ]: # Minimize the squared errors
cost = tf.reduce_sum(tf.pow(y_pred - y, 2)) / (n_samples)

# Define the optimizer
optimizer = tf.train.AdamOptimizer(learning_rate).minimize(cost) #Gradient descent
```

# Linear Regression Example

```
In [6]: # Initializing the variables
init = tf.initialize_all_variables()
# Launch the graph
with tf.Session() as sess:
    sess.run(init)

    # Fit all training data
    for epoch in range(training_epochs):
        sess.run(optimizer, feed_dict={X: train_X, y: train_Y})

        #Display logs per epoch step
        if epoch % display_step == 0:
            print "Epoch:", '%04d' % (epoch+1), "cost=", \
                "{:.9f}".format(sess.run(cost, feed_dict={X: train_X, y: train_Y})), \
                "W=", sess.run(W), "b=", sess.run(b)

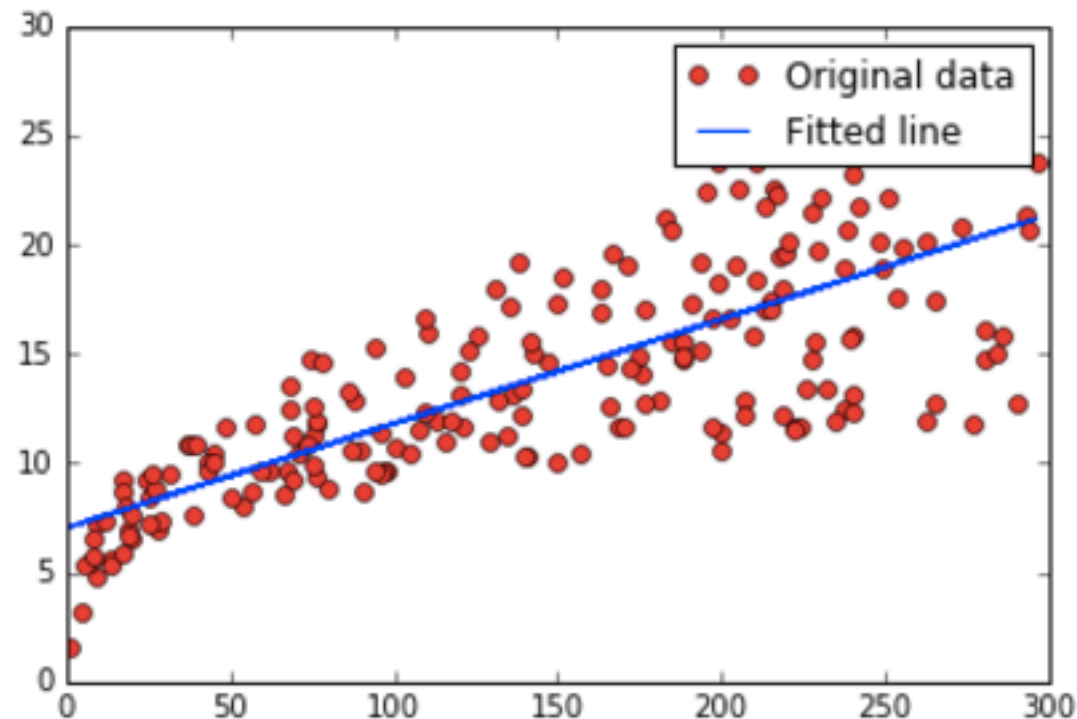
    print "Optimization Finished!"
    print "cost=", sess.run(cost, feed_dict={X: train_X, y: train_Y}), \
        "W=", sess.run(W), "b=", sess.run(b)

    #Graphic display
    plt.plot(train_X, train_Y, 'ro', label='Original data')
    plt.plot(train_X, sess.run(W) * train_X + sess.run(b), label='Fitted line')
    plt.legend()
    plt.show()
```

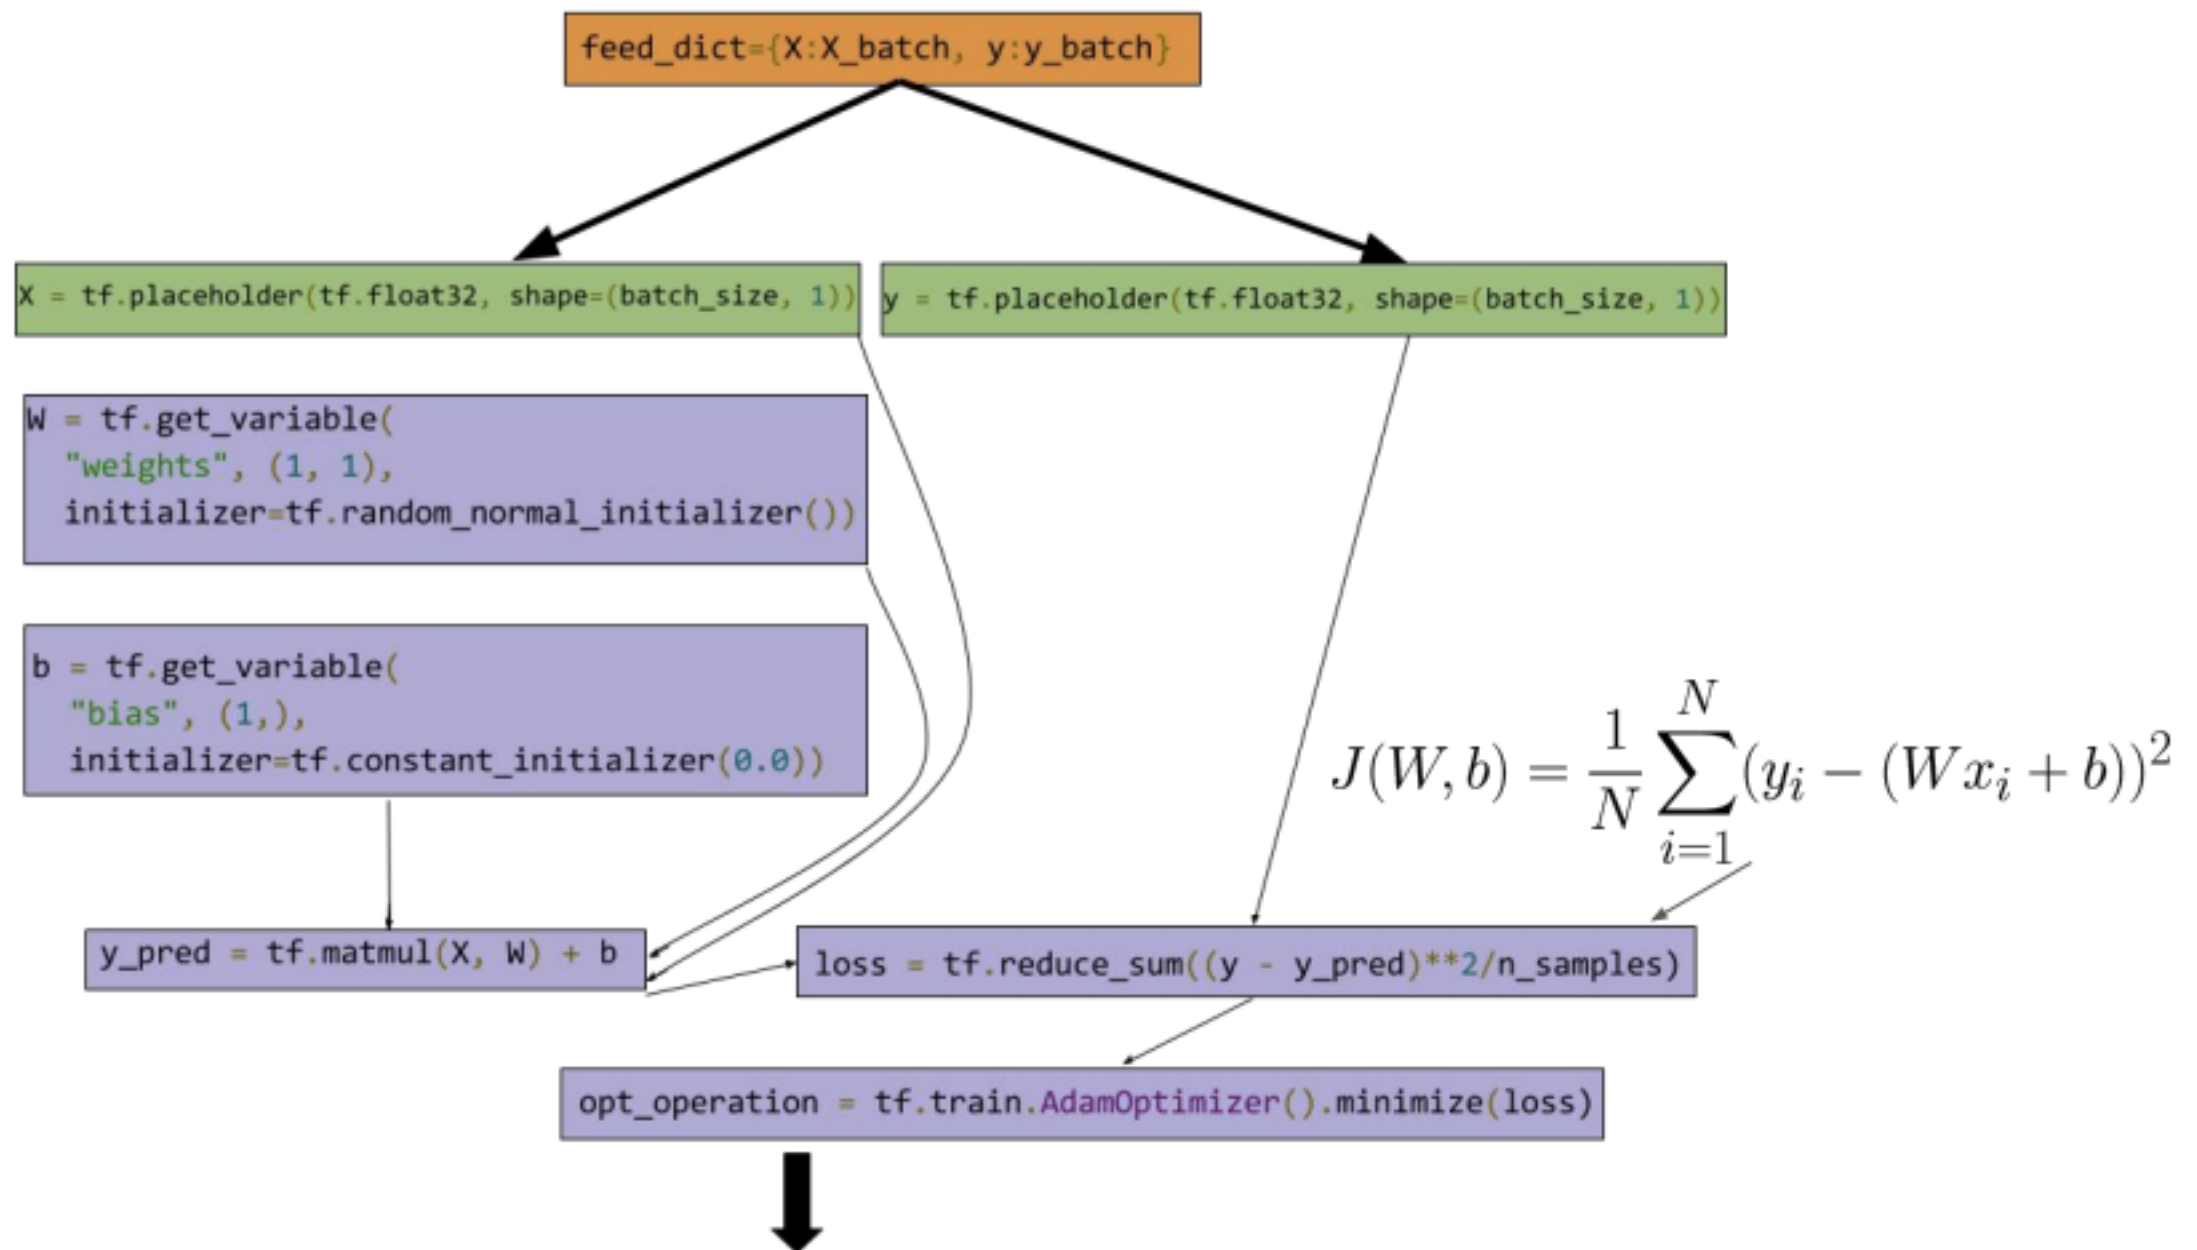


# Linear Regression Example

```
Epoch: 0001 cost= 11.784294128 W= 0.085883 b= -0.113122
Epoch: 0201 cost= 5.516845703 W= 0.0548285 b= 5.59834
Epoch: 0401 cost= 5.256753445 W= 0.0478319 b= 6.97454
Epoch: 0601 cost= 5.256326675 W= 0.0475391 b= 7.03211
Epoch: 0801 cost= 5.256326675 W= 0.0475367 b= 7.03259
Epoch: 1001 cost= 5.256326675 W= 0.0475367 b= 7.03259
Epoch: 1201 cost= 5.256326675 W= 0.0475367 b= 7.03259
Epoch: 1401 cost= 5.256327629 W= 0.0475367 b= 7.03259
Epoch: 1601 cost= 5.256326675 W= 0.0475367 b= 7.03259
Epoch: 1801 cost= 5.256326675 W= 0.0475367 b= 7.03259
Optimization Finished!
cost= 5.25679 W= 0.047714 b= 7.03275
```



# Linear Regression Example



# Which data is stored ?

```
In [6]: # Initializing the variables
init = tf.initialize_all_variables()
# Launch the graph
with tf.Session() as sess:
    sess.run(init)

    # Fit all training data
    for epoch in range(training_epochs):
        sess.run(optimizer, feed_dict={X: train_X, y: train_Y})

        #Display logs per epoch step
        if epoch % display_step == 0:
            print "Epoch:", '%04d' % (epoch+1), "cost=", \
                "{:.9f}".format(sess.run(cost, feed_dict={X: train_X, y: train_Y})), \
                "W=", sess.run(W), "b=", sess.run(b)

    print "Optimization Finished!"
    print "cost=", sess.run(cost, feed_dict={X: train_X, y: train_Y}), \
        "W=", sess.run(W), "b=", sess.run(b)

    #Graphic display
    plt.plot(train_X, train_Y, 'ro', label='Original data')
    plt.plot(train_X, sess.run(W) * train_X + sess.run(b), label='Fitted line')
    plt.legend()
    plt.show()
```

Can we obtain **W** and **b** outside the session?

# Which data is stored ?

```
In [9]: # Initializing the variables
init = tf.initialize_all_variables()
# Launch the graph
with tf.Session() as sess:
    sess.run(init)

    # Fit all training data
    for epoch in range(training_epochs):
        sess.run(optimizer, feed_dict={X: train_X, y: train_Y})

        #Display logs per epoch step
        if epoch % display_step == 0:
            print "Epoch:", '%04d' % (epoch+1), "cost=", \
                  "{:.9f}".format(sess.run(cost, feed_dict={X: train_X, y: train_Y})), \
                  "W=", sess.run(W), "b=", sess.run(b)

    print "Optimization Finished!"
    print "cost=", sess.run(cost, feed_dict={X: train_X, y: train_Y}), \
          "W=", sess.run(W), "b=", sess.run(b)

    #Graphic display
    w = sess.run(W)
    plt.plot(train_X, train_Y, 'ro', label='Original data')
    plt.plot(train_X, w * train_X + sess.run(b), label='Fitted line')
    plt.legend()
    plt.show()
```

# Input Data

- Most of the times we do not have enough memory to load all data from training set and compute the gradients.
- Let's see an example using a batch

# Input Data: Batch

```
In [2]: #import tensorflow
import tensorflow as tf
import numpy as np

# tf Graph Input
X = tf.placeholder("float", [None, 784]) # mnist data image of shape 28*28=784
y = tf.placeholder("float", [None, 10]) # 0-9 digits recognition => 10 classes

# Create model
# Set model weights
W = tf.Variable(tf.zeros([784, 10]))
b = tf.Variable(tf.zeros([10]))

# Construct model
y_pred = tf.nn.softmax(tf.matmul(X, W)) # Softmax
```



# Input Data: Batch

```
In [5]: # Initializing the variables
init = tf.initialize_all_variables()

# Launch the graph
with tf.Session() as sess:
    sess.run(init)

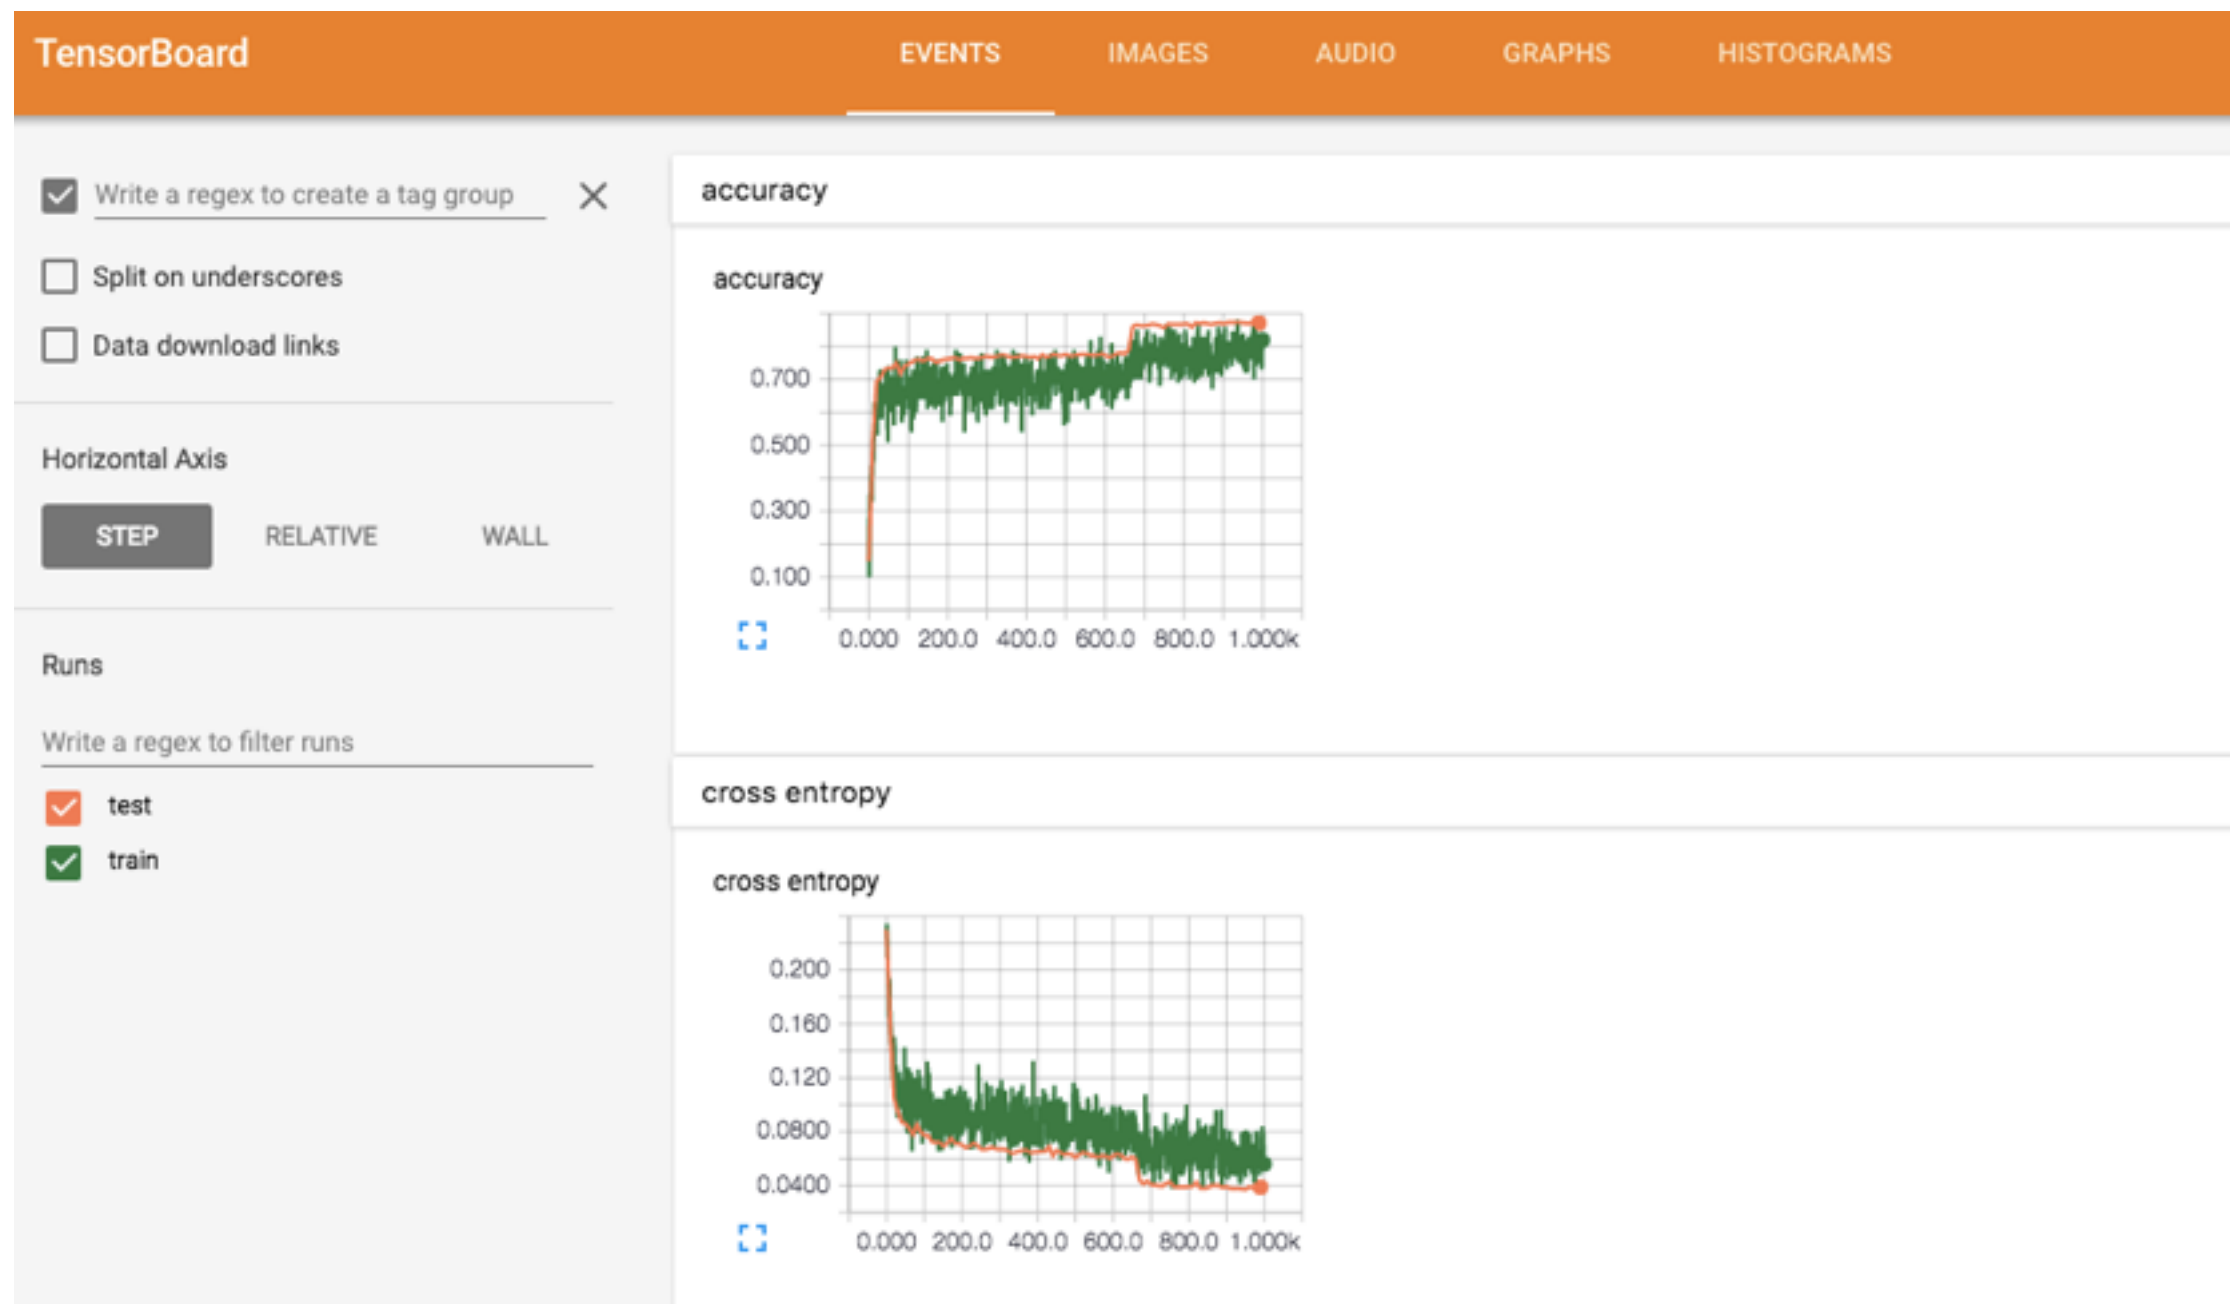
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0.
        total_batch = int(mnist.train.num_examples/batch_size)
        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            # Fit training using batch data
            sess.run(optimizer, feed_dict={X: batch_xs, y: batch_ys})
            # Compute average loss
            avg_cost += sess.run(cost, feed_dict={X: batch_xs, y: batch_ys})/total_batch
        # Display logs per epoch step
        if epoch % display_step == 0:
            print "Epoch:", '%04d' % (epoch+1), "cost=", "{:.9f}".format(avg_cost)

    print "Optimization Finished!"

    # Test model
    correct_prediction = tf.equal(tf.argmax(y_pred, 1), tf.argmax(y, 1))
    # Calculate accuracy
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
    print "Accuracy:", accuracy.eval({X: mnist.test.images, y: mnist.test.labels})
```



# TensorBoard



<https://www.tensorflow.org/tensorboard/index.html>