



UNIVERSITAT DE  
BARCELONA



DATA SCIENCE@UB



Deep Learning From Scratch

# Recurrent Neural Networks

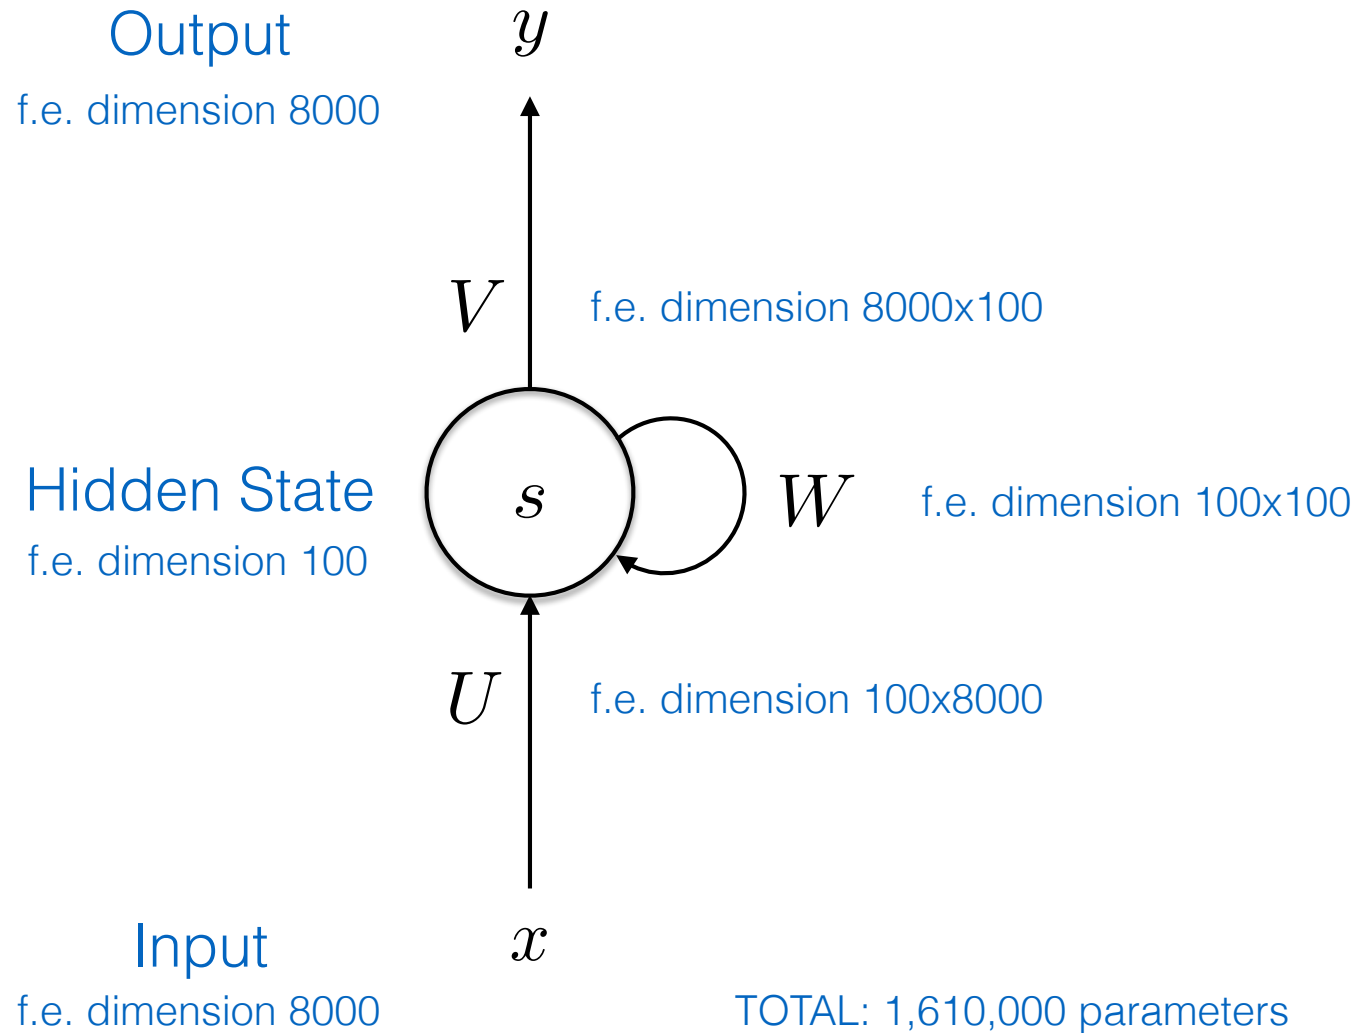
Jordi Vitrià

Classical neural networks, including convolutional ones, suffer from two severe limitations:

- They only accept a fixed-sized vector as input and produce a fixed-sized vector as output.
- They do not consider the sequential nature of some data (language, video frames, time series, etc.)

**Recurrent neural networks** overcome these limitations by allowing to operate over sequences of vectors (in the input, in the output, or both).

# Vanilla Recurrent Neural Network

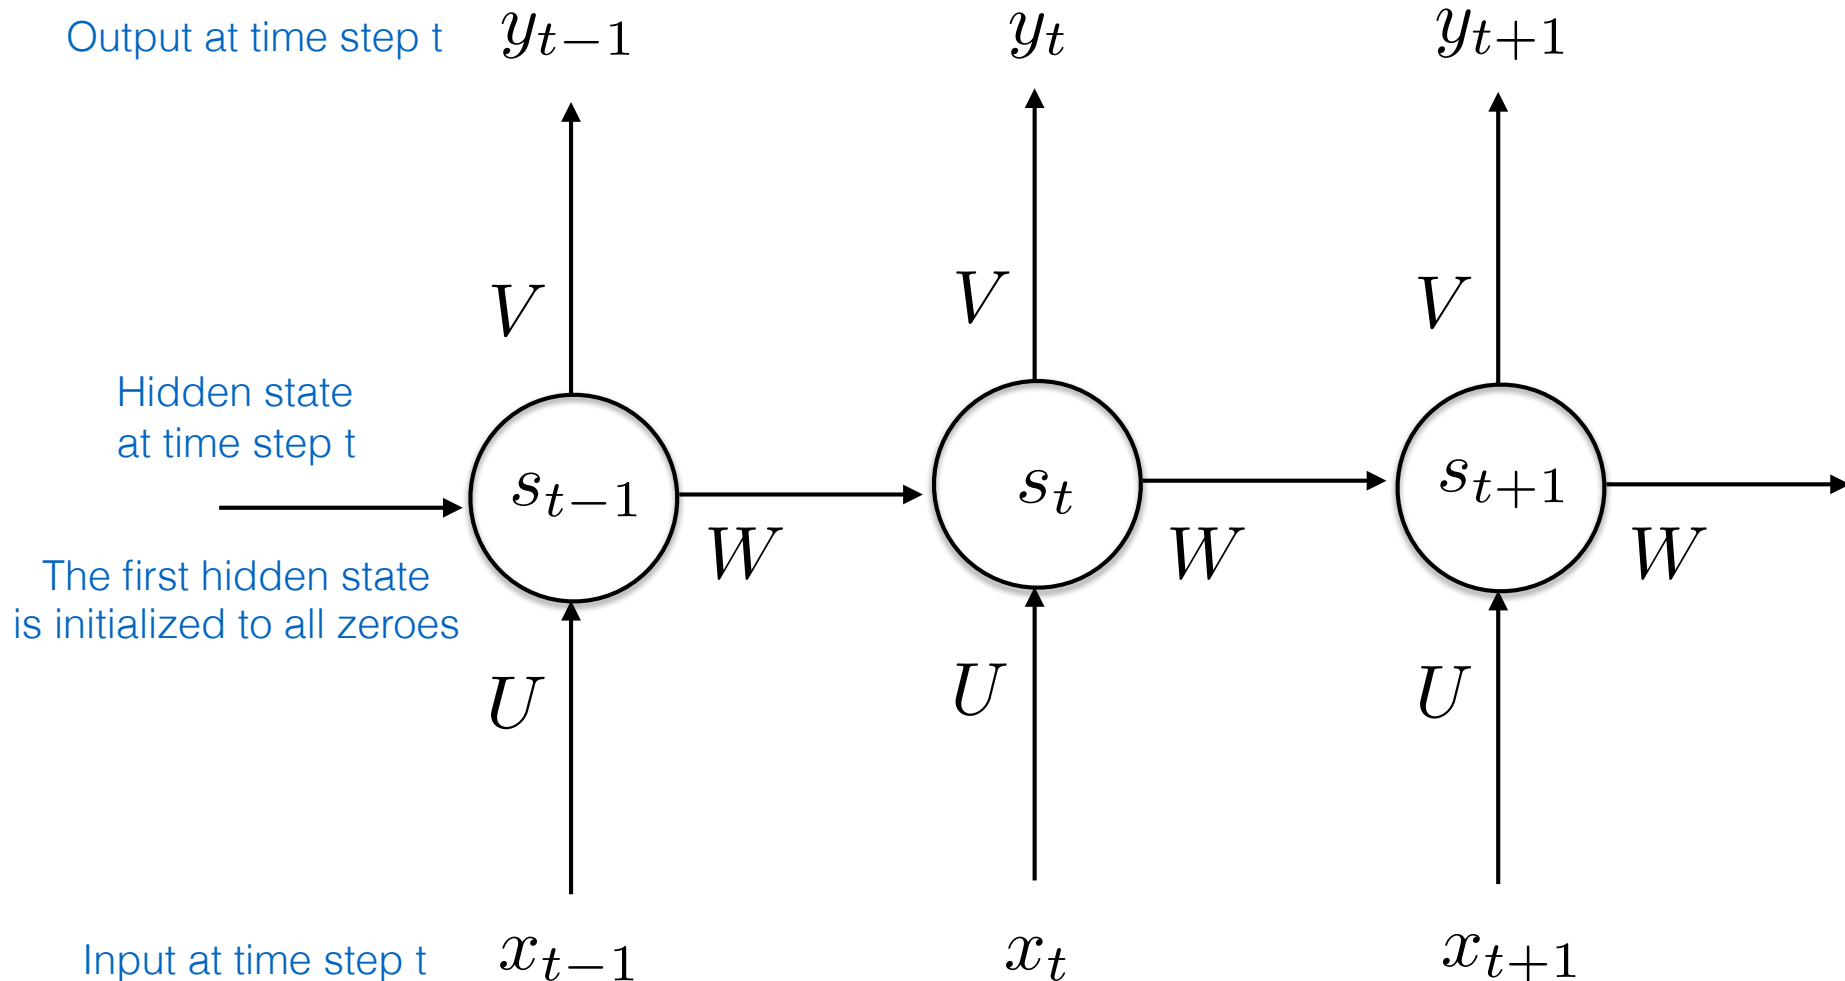


# Unfolding in time of a RNN

By unrolling we mean that we write out the network for the complete sequence.

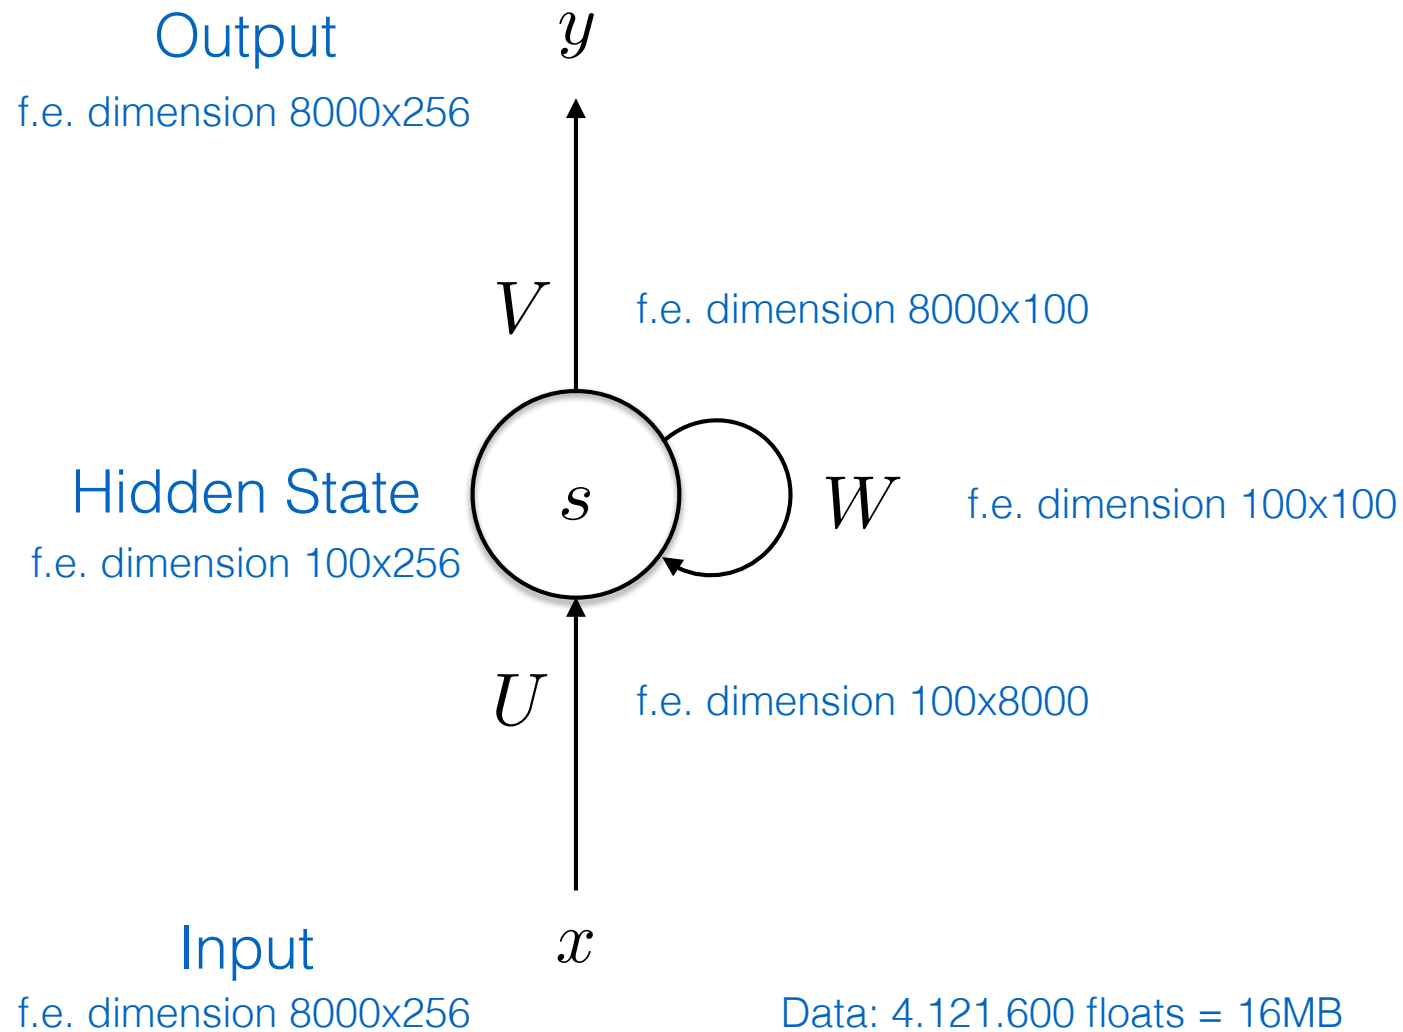
Basic equations of the RNN

$$s_t = \tanh(Ux_t + Ws_{t-1})$$
$$y_t = \text{softmax}(Vs_t)$$



# Vanilla Recurrent Neural Network

minibatch version



$$s_t$$

- We can think of the **hidden state** as a memory of the network that captures information about the previous steps.

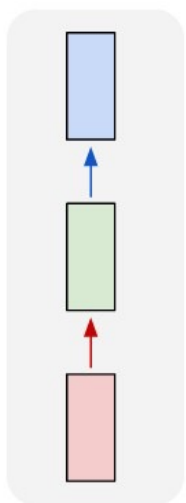
$$U, V, W$$

- The RNN **shares the parameters** across all time steps.

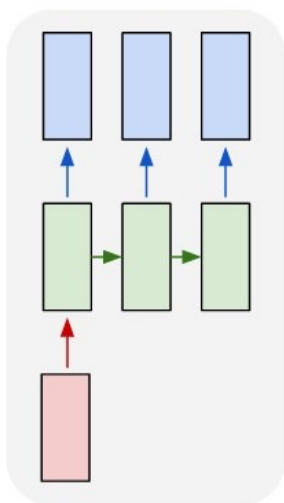
$$y_t$$

- It is not necessary to have **outputs** at each time step.

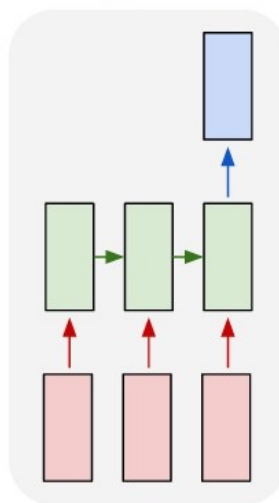
one to one



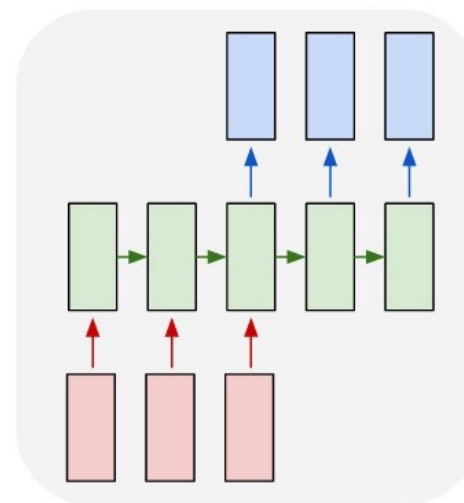
one to many



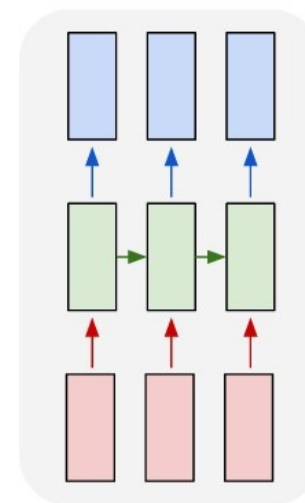
many to one



many to many



many to many



Source: <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

RNN have shown success in:

- Language modeling and generation.
- Machine Translation.
- Speech Recognition.
- Image Description.
- Question Answering.
- Etc.



# RNN Training

Training a RNN is similar to training a traditional NN, but some modifications.

The main reason is that parameters are shared by all time steps: in order to compute the gradient at  $t=4$ , we need to propagate 3 steps and sum up the gradients.

This is called **Backpropagation through time** (BPTT).

# RNN Computation

```
class RNN:
    # ... self.h is initialized with the zero vector.
    def step(self, x):
        self.h = np.tanh(np.dot(self.W_hh, self.h) + np.dot(self.W_xh, x))
        y = np.dot(self.W_hy, self.h)
        return y
    # ... Matrices are initialized with random numbers.
```

We can go deep by stacking RNN:

```
y1 = rnn1.step(x)
y2 = rnn2.step(y1)
```

# RNN Models

Vanilla RNNs trained with SGD are unstable/difficult to learn. But various **tricks** make our life easier:

- Gating Units
- Gradient Clipping
- Steeper gates
- Better initialization

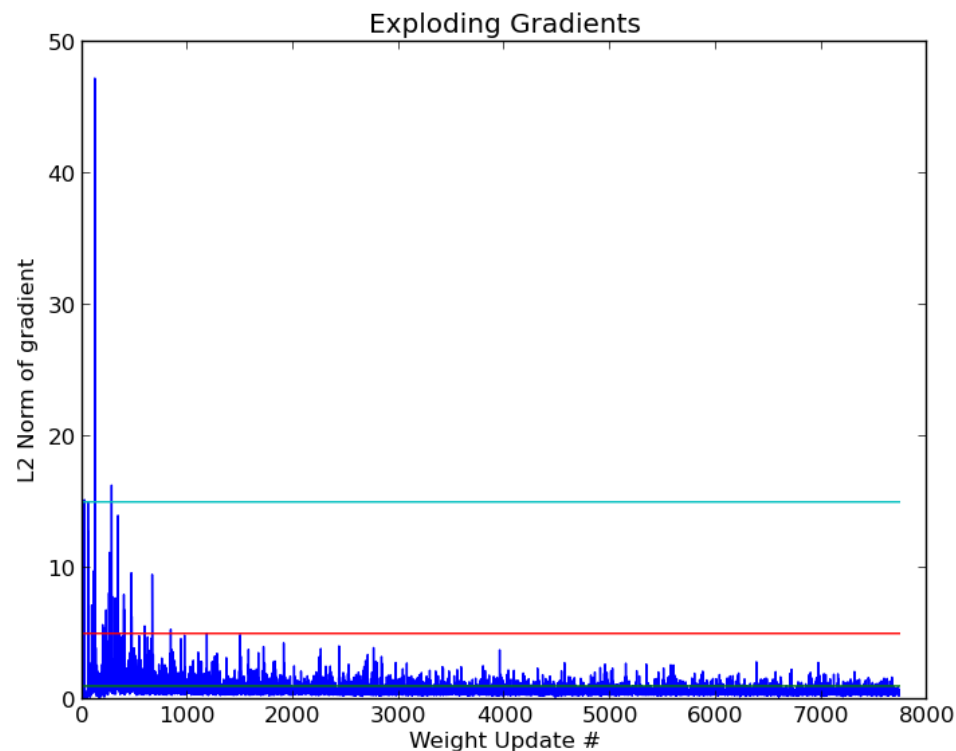
# Gated Units

There are two types of gated RNNs:

- **Gated Recurrent Units** (GRU) by recently introduced K. Cho. GRU is simpler, faster, and optimizes quicker.
- **Long short term memory** (LSTM) by S. Hochreiter and J. Schmidhuber has been around since 1997 and has been used far more. LSTM may be better in the long run due to its greater complexity.

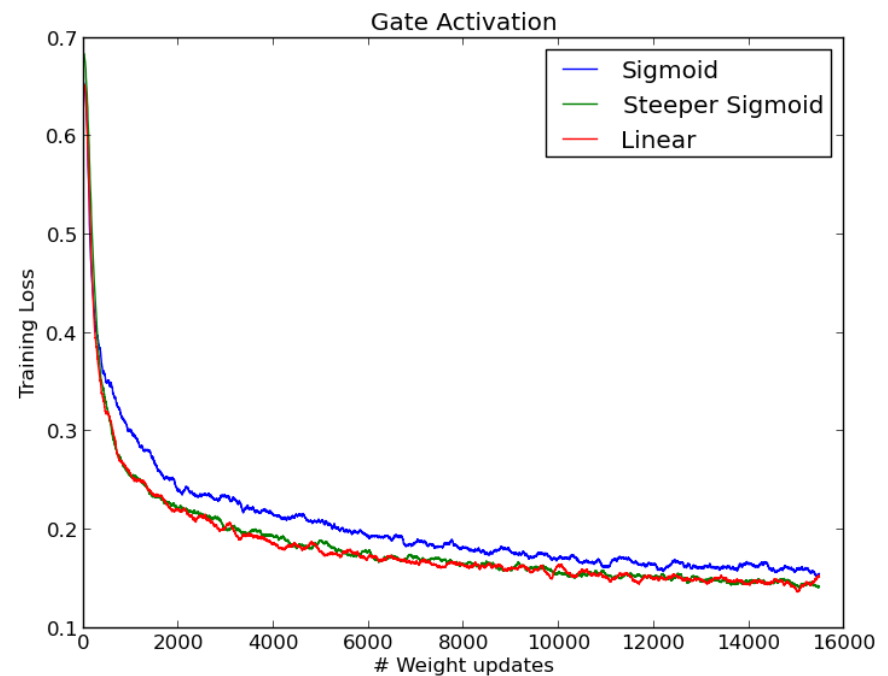
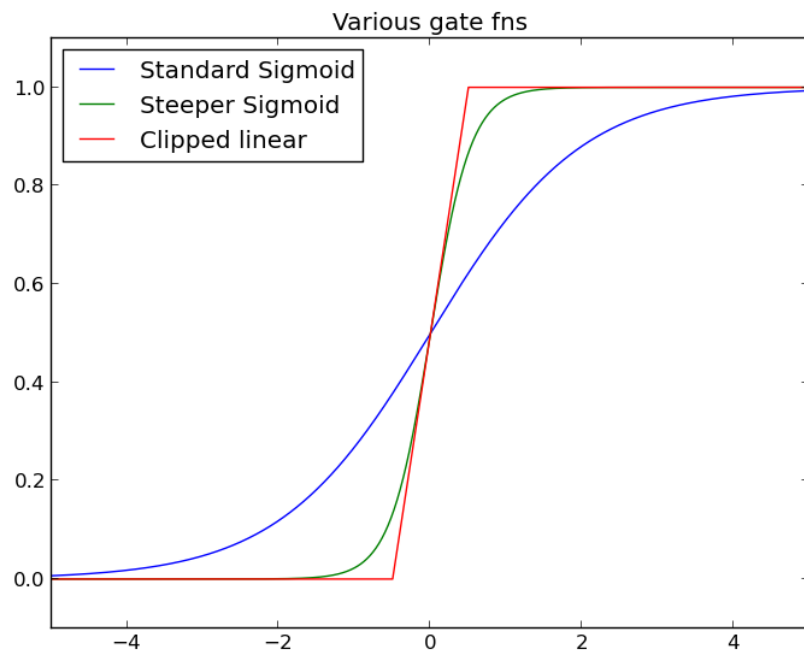
# Exploding gradients

Exploding gradients may be a major problem for traditional RNNs trained with SGD. In 2012, R Pascanu and T. Mikolov proposed clipping the norm of the gradient to alleviate this.



# Steeper Gates

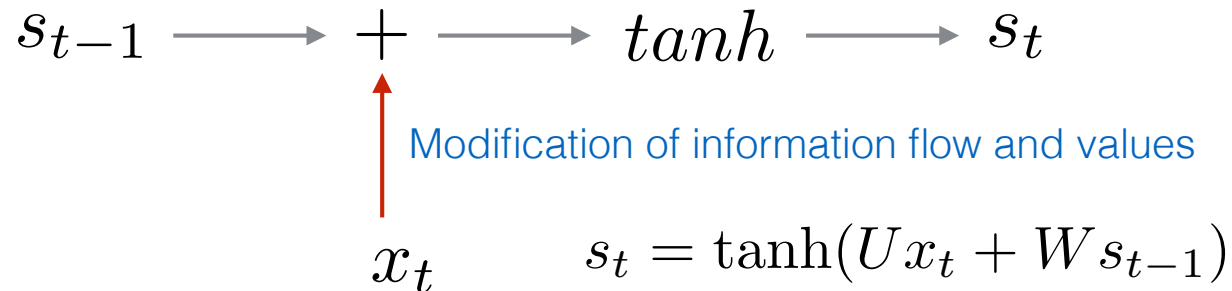
We can make the gates “steeper” so they change more rapidly from “off” to “on” so model learns to use them quicker.



# Better initialization

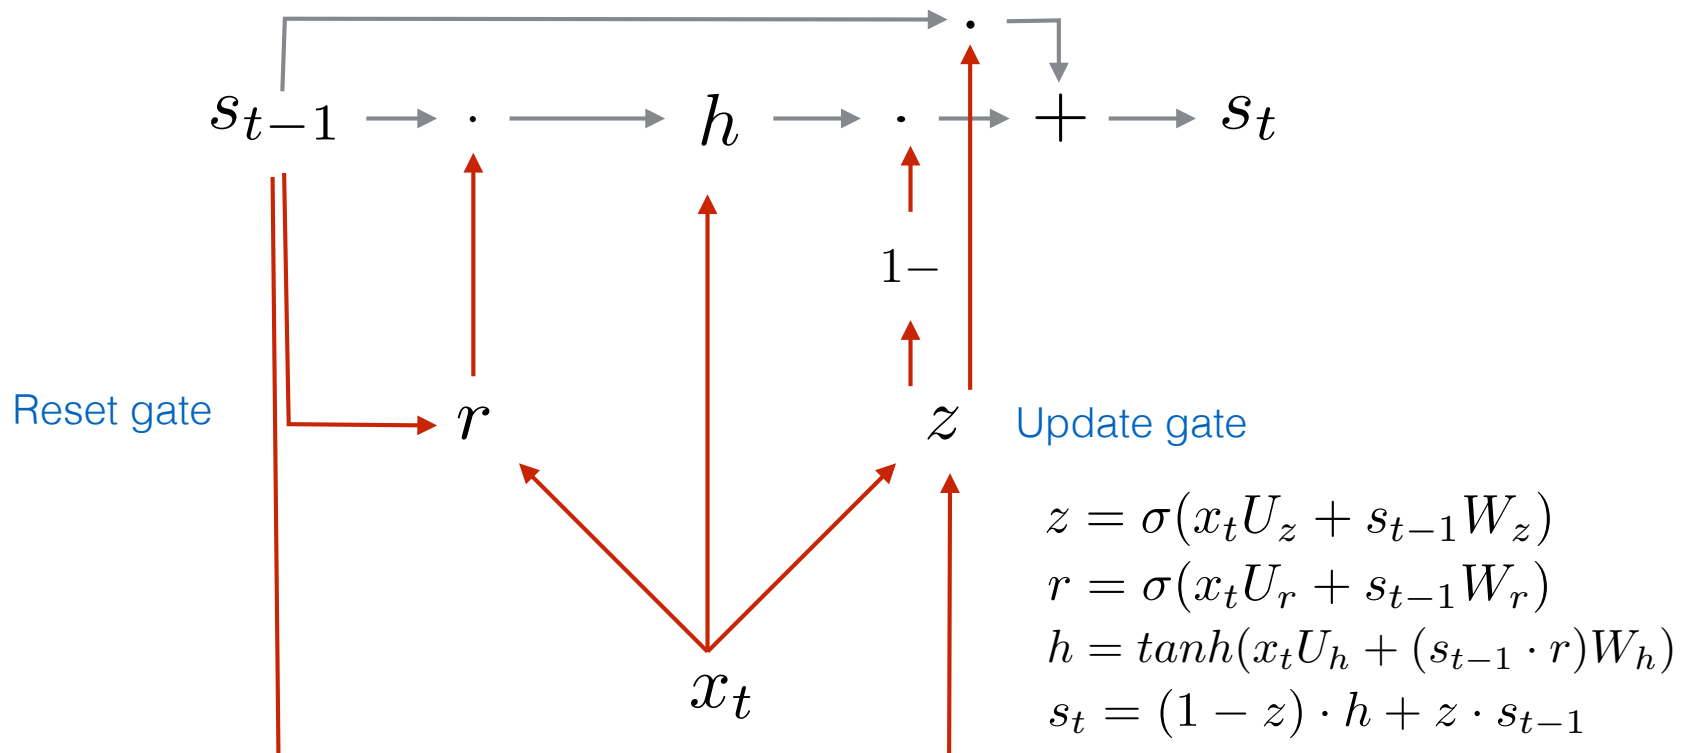
It has been showed that initializing weight matrices with random orthogonal matrices works better than random gaussian (or uniform) matrices.

# Gated Recurrent Unit (GRU)



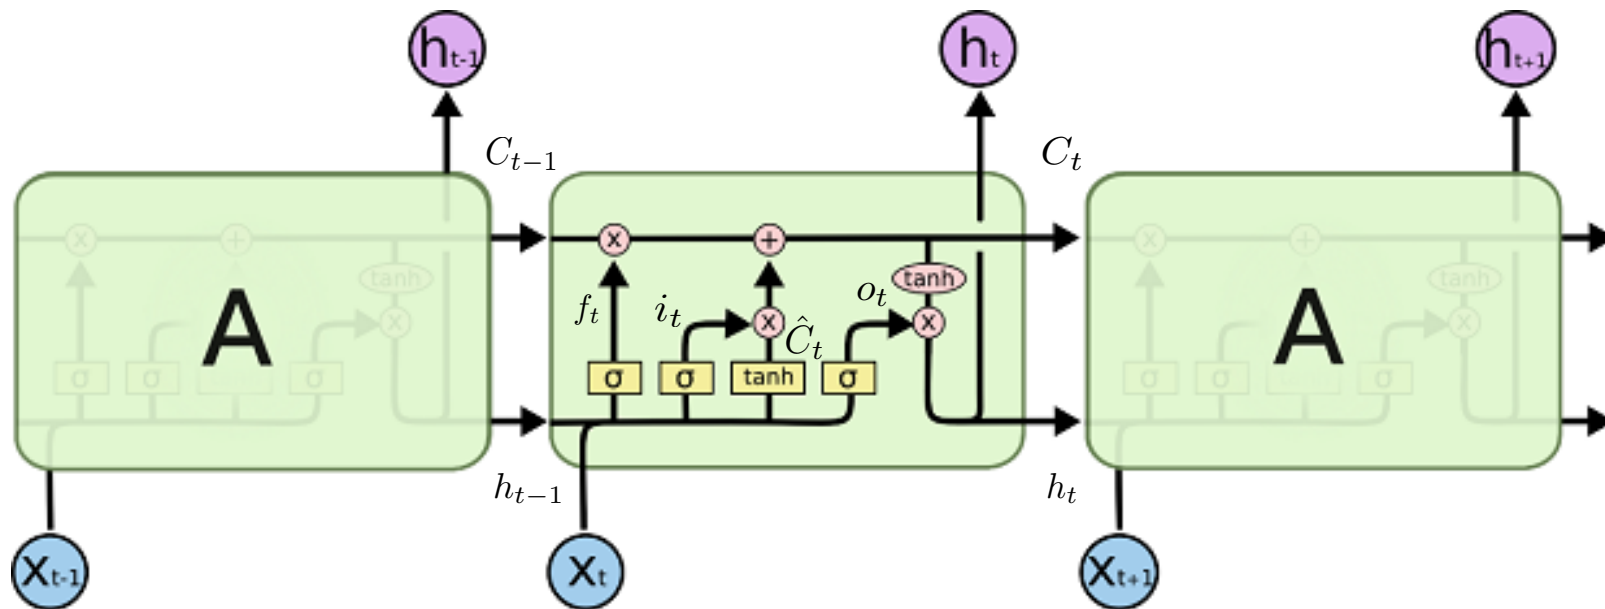
Vanilla RNN

GRU





# Long Short Term Memory Unit (LSTM)



Source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

$$\begin{aligned}f_t &= \sigma(W_f[h_{t-1} \cdot x_t] + b_f) \\i_t &= \sigma(W_i[h_{t-1} \cdot x_t] + b_i) \\\hat{C}_t &= \tanh(W_C[h_{t-1} \cdot x_t] + b_C) \\C_t &= f_t C_{t-1} + i_t \hat{C}_t \\o_t &= \sigma(W_o[h_{t-1} \cdot x_t] + b_o) \\h_t &= o_t \tanh(C_t)\end{aligned}$$

# From text to RNN input

String Input

**The cat sat on the mat.**

Tokenize

**The cat sat on the mat .**

Indexing

**0 1 2 3 0 4 5**

Embedding

**2.5 0.3 -1.2**

**0.2 -3.3 0.7**

**-4.1 1.6 2.8**

**1.1 5.7 -0.2**

**2.5 0.3 -1.2**

**1.4 0.6 -3.9**

**-3.8 1.5 0.1**

# Example: Name Modeling

Let's build a sequential (Name) model with a Recurrent Neural Network. Let's say we have name of  $m$  chars.

A name model allows us to predict the probability of observing the name as:

$$P(c_1 \dots c_m) = \prod_{i=1}^m P(c_i | c_1 \dots c_{i-1})$$

Note that in the equation the probability of each char is conditioned on all previous chars.

# Example: Name Modeling

To train our model we need text to learn from a large dataset of names. Fortunately we don't need any labels to train a language model, just raw text.

I downloaded 52,700 Catalan names from a dataset available on



[http://territori.gencat.cat/ca/01\\_departament/  
11\\_normativa\\_i\\_documentacio/  
03\\_documentacio/02\\_territori\\_i\\_mobilitat/  
cartografia/  
nomenclator\\_oficial\\_de\\_toponimia\\_de\\_catalunya/](http://territori.gencat.cat/ca/01_departament/11_normativa_i_documentacio/03_documentacio/02_territori_i_mobilitat/cartografia/nomenclator_oficial_de_toponimia_de_catalunya/)

# Example: Name Modeling

## Results

Alzinetes, torrent de les	Regueret, lo
Alzinetes, vall de les	Regueret, lo
<b>Alzinó, Mas d'</b>	<b>Regueró</b>
Alzinosa, collada de l'	Reguerols, els
Alzinosa, font de l'	Reguerons, els
Benavent, roc de	Vallverdú, Mas de
Benaviure, Cal	Vallverdú, serrat de
<b>Benca</b>	<b>Vallvicamanyà</b>
Bendiners, pla de	Vallvidrera
Benedi, roc del	Vallvidrera, riera de
Fiola, la	Terraubella, Corral de
Fiola, puig de la	Terraubes
<b>Fiper, Granja del</b>	<b>Terravanca</b>
Firassa, Finca	Terrer Nou, Can
Firell	Terrer Roig, lo